

THE SENTINEL REVIEW INTEL DASHBOARD

White Paper — May 2026

Executive Summary

The Sentinel Review Intel Dashboard is a free, publicly accessible open-source intelligence (OSINT) platform that aggregates, verifies, and visualizes conflict events in near real-time. By combining automated ingestion from over 35 curated intelligence sources with AI-assisted event extraction and deterministic confidence scoring, the dashboard provides policymakers, analysts, journalists, and defense professionals with a transparent, structured picture of active conflict theaters.

The dashboard is live at dashboard.thesentinelreview.com and is fully integrated with The Sentinel Review publication at thesentinelreview.com.

1. What the Dashboard Is

The Sentinel Review Intel Dashboard is a structured OSINT aggregation and verification platform. It is not a primary source — it is a synthesis layer that pulls raw reporting from wire services, investigative accounts, regional press, and open Telegram/X channels, processes each post through an AI extraction pipeline, scores the resulting events by confidence, and presents them on an interactive map with daily AI-generated briefings.

The dashboard is explicitly positioned as a situational awareness tool, not an operational one. It does not support military targeting or tasking functions. Events involving mass casualties, weapons of mass destruction, or cross-border escalation are held from public display until a human editor has reviewed and approved them.

The current v0.1 release covers the Ukraine theater with a focus on the Donetsk, Luhansk, Zaporizhzhia, and Kharkiv oblasts.

2. How It Functions

2.1 Data Ingestion

Every 30 minutes, an automated pipeline fetches new posts from all active sources across four platforms:

- RSS — wire services and established publications (Reuters, AP, AFP, Kyiv Independent, Defense News)
- X (Twitter) — verified OSINT accounts (DefMon3, GeoConfirmed, OSINTtechnical, UAWeapons)
- Telegram — frontline channels (DeepStateUA, Militaryland, Rybar, UkraineWarReport)
- Wire — structured event databases (ACLED, used for backfill and cross-validation)

All raw posts are written to an append-only database log. Nothing is deleted or mutated — every post is preserved for audit purposes.

2.2 AI Event Extraction

Each raw post is passed to Claude Sonnet 4.6 with a structured system prompt authored for an OSINT analyst role. The model uses structured JSON tool use, which eliminates free-text hallucination. For each post, the model determines:

- Whether a discrete conflict event is present
- Event type: strike, clash, or movement
- The time the event occurred (extracted from text or defaulted to post timestamp)
- Location name, oblast, and estimated WGS84 coordinates
- The acting force if identifiable
- A factual 1–2 sentence description with no editorializing
- Geolocation signals: geolocated footage, official acknowledgment, matching press, coordinates, visible landmarks
- Whether the event is high-impact (mass casualties ≥ 10 , WMD signal, cross-border escalation)

Every LLM call is logged to an audit table with the full prompt, response, and token counts.

2.3 Deduplication

Before an event is scored, the pipeline checks for near-duplicates. Events reported within a 5 km radius and 6-hour window of an existing event are treated as corroborating reports of the same incident, not new events. This prevents the same strike from appearing multiple times because three channels reported it independently.

2.4 Confidence Scoring

Confidence is assigned by deterministic rules — no probabilistic model:

Verified	≥2 independent sources from different platforms + at least one of: geolocated footage, official acknowledgment, or matching press coverage
Partial	≥2 sources on the same platform, or 1 Tier 1 source + 1 corroborating signal
Unconfirmed	Single source, or multiple sources sharing a common origin

Source trust tiers weight this scoring:

- Tier 1 — Verified investigative accounts (DefMon3, GeoConfirmed), wire services (Reuters, AP)
- Tier 2 — Established milblogs and regional press (ISW, Militaryland, Kyiv Independent)
- Tier 3 — Anonymous channels, state-affiliated media (Rybar, TASS) — corroboration only, never sufficient alone for Verified status

2.5 Human Review Gate

High-impact events are held from publication automatically and routed to a human editor queue. Editors can upgrade or downgrade confidence, add notes, and publish or reject. No event appears on the public dashboard until it is approved by a human reviewer.

All changes to events — including corrections — are preserved in a public change history. Nothing is silently edited.

2.6 Daily Briefings

Every morning at 06:00 UTC, Claude Opus 4.7 generates a 2–3 paragraph briefing covering the past 24 hours in the theater. The briefing summarizes activity volume, geography, and shifts vs. the 7-day baseline; describes key verified and partial events; and closes with a watch item for the next 24 hours.

All briefings are marked "AI DRAFT" until a human editor reviews and publishes them. The model is instructed to distinguish verified events from unconfirmed ones in prose.

2.7 Dashboard Interface

The dashboard presents all of this through a Next.js web application with the following components:

- Interactive map — geolocated conflict events color-coded by type (red = strikes, amber = clashes, blue = movements), rendered via MapLibre GL
- At a Glance panel — event counts, strike counts, verified percentage, 7-day trend
- Intensity chart — 7-day bar chart highlighting anomalous days against the daily average baseline

- Active alerts — the three most recent events with confidence badges, linking to full event detail pages
- Top sources — ranked by volume and 30-day verification rate
- Daily briefing — latest published briefing with embedded event references
- Supporting pages — /sources (source directory with reliability metrics), /methodology (scoring rules), and individual event and briefing pages

3. Benefits

3.1 Speed Without Sacrificing Rigor

Traditional conflict monitoring requires analysts to read dozens of channels manually. The pipeline compresses this into minutes: a post published on Telegram at 14:30 UTC can be on the public map by 14:35 UTC, with confidence already assessed and sources attributed.

3.2 Transparent Confidence

Every event on the dashboard carries its confidence level, the specific posts that contributed to it, and the relationship between sources (primary, corroborating, contradicting). Users can drill into any event and read the original posts.

3.3 Source Accountability

The 30-day rolling verified rate for each source is publicly visible on the /sources page. Sources that consistently produce unconfirmable claims are visible as such, giving users a basis for evaluating reliability that does not require institutional knowledge of the OSINT space.

3.4 Editorial Independence from LLMs

The AI does extraction and drafting — it does not publish. High-impact events require human sign-off. Briefings require human review. Corrections are permanent and public. The editorial layer is human; the AI layer is a labor-saving instrument underneath it.

3.5 Cost Efficiency

The entire pipeline runs on approximately \$50–200/month at current ingestion volume (35 sources, ~100–200 events/day). The primary cost driver is Claude API calls for extraction — orders of magnitude cheaper than maintaining an analyst team for equivalent coverage.

3.6 Open and Embeddable

Briefings are embeddable. The dashboard is public and requires no account. The methodology is fully documented. There are no paywalls on the situational awareness layer.

4. Use Cases

Policy and government staff — Rapid situational awareness without wading through raw Telegram channels; confidence scoring surfaces what is verified vs. what is noise.

Journalists and editors — Source discovery and cross-referencing; event detail pages show exactly which accounts reported what and when, reducing verification legwork.

Think tanks and research institutions — Structured event data for quantitative analysis; the database schema is designed for export and longitudinal study.

Defense and security professionals — Theater-level activity summaries at a glance; the intensity chart flags anomalous days against the 7-day baseline.

General public — Accessible, free, plain-language briefings with transparent sourcing for anyone following the conflict.

5. Design Principles

Transparency over authority. Every confidence level, every source, every AI call is logged and exposed. The dashboard does not claim to be definitive — it claims to be honest about what it knows and how it knows it.

Editorial integrity over speed. High-impact events are held. Briefings require review. Corrections are permanent. Speed is optimized within these constraints, not at their expense.

No operational use. The platform is explicitly a situational awareness tool. It does not support targeting, military planning, or any kinetic use. This is documented publicly on the /methodology and /about pages.

Deterministic confidence. Confidence scores follow explicit, published rules — not black-box model outputs. This makes the system auditable and the scores explainable to any reader.

6. Technical Stack

Layer	Technology
Web application	Next.js 15+, TypeScript, MapLibre GL
Hosting	Vercel
Ingest pipeline	Python 3.11, psycopg3
Job queue	PostgreSQL (no Redis required)
Database	PostgreSQL 15+ with PostGIS
Managed DB	Neon Postgres
AI extraction	Claude Sonnet 4.6 (Anthropic)
AI briefings	Claude Opus 4.7 (Anthropic)
Workflow automation	GitHub Actions (ingest every 30 min, briefing daily at 06:00 UTC)
DNS / CDN	Cloudflare

7. Roadmap

- v0.1 (current) — Ukraine theater; RSS + X + Telegram ingestion; map, alerts, briefings, source directory; GitHub Actions pipeline
- v0.2 — Admin briefing review interface; Telegram worker activation; map event clustering; time scrubber for historical playback; embed modal for briefings
- v0.3 — Second theater (TBD based on editorial priority); ACLED backfill integration; batch API for cost reduction; public event API